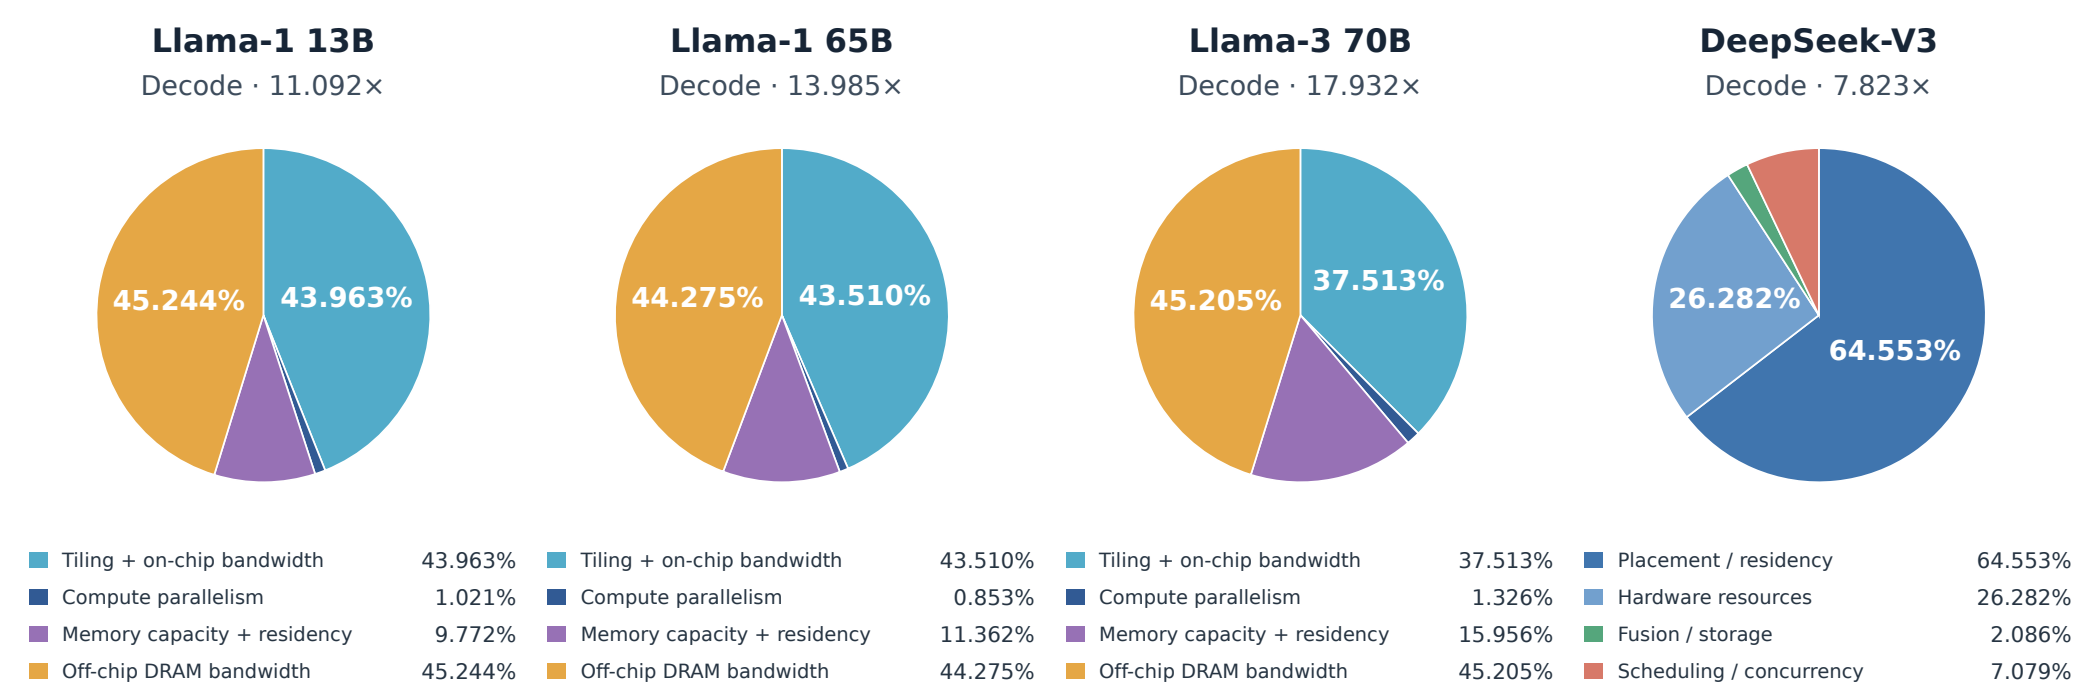
Llama-1 13B
Decode · 11.092×
45.244%
43.963%
Llama-1 65B
Decode · 13.985×
44.275%
43.510%
Llama-3 70B
Decode · 17.932×
45.205%
37.513%
DeepSeek-V3
Decode · 7.823×
26.282%
64.553%
Tiling + on-chip bandwidth 43.963%
Compute parallelism 1.021%
Memory capacity + residency 9.772%
Off-chip DRAM bandwidth 45.244%
Tiling + on-chip bandwidth 43.510%
Compute parallelism 0.853%
Memory capacity + residency 11.362%
Off-chip DRAM bandwidth 44.275%
Tiling + on-chip bandwidth 37.513%
Compute parallelism 1.326%
Memory capacity + residency 15.956%
Off-chip DRAM bandwidth 45.205%
Placement / residency 64.553%
Hardware resources 26.282%
Fusion / storage 2.086%
Scheduling / concurrency 7.079%
Grouped feature contributions · 100% of modeled speedup gain above 1×