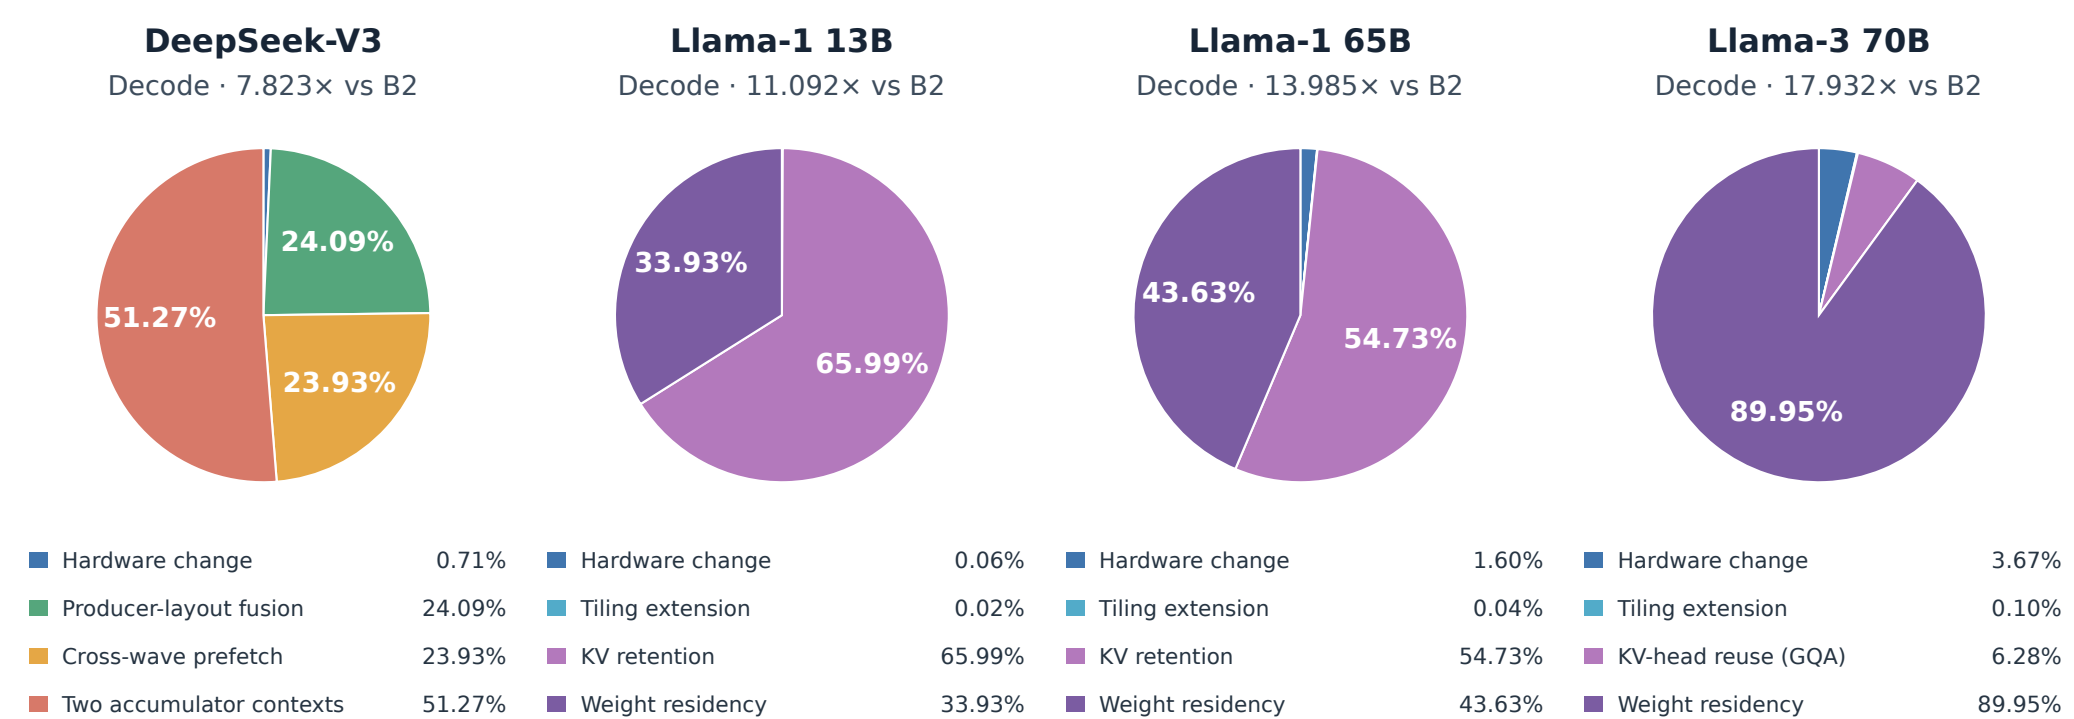
DeepSeek-V3
Decode · 7.823× vs B2
24.09%
23.93%
51.27%
Hardware change 0.71%
Producer-layout fusion 24.09%
Cross-wave prefetch 23.93%
Two accumulator contexts 51.27%
Llama-1 13B
Decode · 11.092× vs B2
33.93%
65.99%
Hardware change 0.06%
Tiling extension 0.02%
KV retention 65.99%
Weight residency 33.93%
Llama-1 65B
Decode · 13.985× vs B2
43.63%
54.73%
Hardware change 1.60%
Tiling extension 0.04%
KV retention 54.73%
Weight residency 43.63%
Llama-3 70B
Decode · 17.932× vs B2
89.95%
Hardware change 3.67%
Tiling extension 0.10%
KV-head reuse (GQA) 6.28%
Weight residency 89.95%
Visual estimates · feature weights normalized to 100% · overlapping effects (see caption)